

You might notice an output asking you to look at a log file to identify what went wrong. This is because you might not be using a **Graphical User Interface (GUI)** for Linux. The error is simply stating that there was no GUI identified and that it was skipped. The output should appear as follows – in which case it can be safely ignored:

```
VirtualBox Guest Additions: Look at /var/log/vboxadd-setup.log  
to find out what went wrong
```

Once we have completed that (it takes a few minutes), let's reboot the Elastic VM:

```
$ sudo reboot
```

Now that we've fully prepared the Elastic VM, we can proceed with installing Elasticsearch.

## Summary

In this chapter, we explored the architecture that you'll use for your lab environment. Additionally, we built the Elastic VM and performed the basic preparatory steps needed to install the Elastic Stack.

In the next chapter, we will install the various components of the Elastic Stack, the victim VM, and ingest threat data into the Elastic Stack.

## Questions

As we conclude, here is a list of questions for you to test your knowledge regarding this chapter's material. You will find the answers in the *Assessments* section of the *Appendix*:

1. What is the name of the machine that the VMs reside on?
  - a. Guest
  - b. Host
  - c. Virtual machine
  - d. Container

2. What does VirtualBox function as?
  - a. A hypervisor
  - b. A guest
  - c. An operating system
  - d. The Elastic Stack
3. What is the DNF program used for?
  - a. Starting services
  - b. Configuring network interfaces
  - c. Installing software on Linux
  - d. Controlling user accounts
4. Prior to installing VirtualBox Guest Additions, what must you do?
  - a. Remove the Linux ISO from VirtualBox.
  - b. Reboot the system.
  - c. Install a web browser.
  - d. Disable the network interface.
5. Which command updates the Linux system?
  - a. `sudo dnf upgrade`
  - b. `sudo dnf patch`
  - c. `sudo apt-get patch`
  - d. `sudo dnf update`

# 5

# Building Your Hunting Lab – Part 2

Now that we've discussed the architecture and built our Elastic Virtual Machine (VM), let's continue with installing and configuring the components of the Elastic Stack and our victim VM and ingest some threat information into the stack.

Keeping with the process in previous chapters, we'll use this chapter to build and the next chapter (*Chapter 6, Data Collection with Beats and Elastic Agent*) to install and configure the host components on the victim machine.

In this chapter, we'll go through the following topics:

- Installing and configuring Elasticsearch
- Installing Elastic Agent
- Installing and configuring Kibana
- Enabling the detection engine and Fleet
- Building a victim machine
- Filebeat Threat Intel module

## Technical requirements

In this chapter, you will need to have access to the following:

- VirtualBox (or any hypervisor) with at least 12 GB of RAM, six CPU cores, and 70 GB HDD available to VM guests.
- A Unix-like operating system (macOS, Linux, and so on) is strongly recommended.
- A text editor that will not add formatting (Sublime Text, Notepad++, Atom, vi/vim, Emacs, nano, and so on).
- Access to a command-line interface.
- The archive program `tar`.
- A modern web browser with a UI.
- A package manager is recommended, but not required.
- macOS Homebrew – <https://brew.sh>.
- Ubuntu APT – included in Ubuntu-like systems.
- RHEL/CentOS/Fedora yum or DNF – included in RHEL-like systems.
- Windows Chocolatey – <https://chocolatey.org/install>.

### Important note

We'll be building a sandbox to eventually detonate malware for dynamic analysis. It is essential to remember that while we're taking steps to ensure our host is staying secure, we are going to be detonating malicious software that while extremely rare could have the potential to escape a hypervisor. Treat the malware and packet captures carefully to ensure there is not an accidental infection, using segmented infrastructure if possible.

The code for the examples in this chapter can be found at the following GitHub link:  
[https://github.com/PacktPublishing/Threat-Hunting-with-Elastic-Stack/tree/main/chapter\\_5\\_building\\_your\\_hunting\\_lab\\_part\\_2](https://github.com/PacktPublishing/Threat-Hunting-with-Elastic-Stack/tree/main/chapter_5_building_your_hunting_lab_part_2).

Check out the following video to see the Code in Action:  
<https://bit.ly/3wHF2te>

# Installing and configuring Elasticsearch

As we move forward in the chapter (and beyond), we'll not need to repeat these steps as Kibana, Fleet, and the detection engine all reside on the same guest.

## Adding the Elastic repository

As discussed previously, using a package manager is much cleaner and easier than simply running binaries as we did in some examples in the previous chapter.

Once again, we'll be using yum or DNF as our package manager, but first, we need to add the Elastic repositories.

We'll use nano as our text editor (because it's a bit easier), but feel free to use vim or the like if you're more comfortable (or any other text editor).

Let's create the `elastic.repo` file in the `/etc/yum.repos.d` directory:

```
$ sudo nano /etc/yum.repos.d/elastic.repo
[elastic]
name=Elastic repository for 7.x packages
baseurl=https://artifacts.elastic.co/packages/7.x/yum
gpgcheck=1
gpgkey=https://artifacts.elastic.co/GPG-KEY-elasticsearch
enabled=1
autorefresh=1
type=rpm-md
```

Next, let's import the Elasticsearch signing key to validate the installations:

```
$ sudo rpm --import https://artifacts.elastic.co/GPG-KEY-
elasticsearch
```

## Installing Elasticsearch

Now that we have prepped the system to install Elasticsearch, we can run the installation using yum or DNF:

```
$ sudo dnf install elasticsearch
```

Once the installation is complete, let's do a quick functions check before we move on to deploying the required security configuration.

Let's set Elasticsearch to start on boot (`systemctl enable`) and start it (`systemctl start`):

```
$ sudo systemctl enable elasticsearch
$ sudo systemctl start elasticsearch
```

As we did in the previous chapter, let's hit the Elasticsearch API with the cURL program to validate that it's working:

```
$ curl localhost:9200

{
  "name" : "elastic-packetpub.local",
  "cluster_name" : "elasticsearch",
  "cluster_uuid" : "_yiyiDYdQ620Y6FjODu1JQ",
  "version" : {
    "number" : "7.11.1",
    "build_flavor" : "default",
    "build_type" : "rpm",
    "build_hash" : "ff17057114c2199c9c1bbecc727003a907c0db7a",
    "build_date" : "2021-02-15T13:44:09.394032Z",
    "build_snapshot" : false,
    "lucene_version" : "8.7.0",
    "minimum_wire_compatibility_version" : "6.8.0",
    "minimum_index_compatibility_version" : "6.0.0-beta1"
  },
  "tagline" : "You Know, for Search"
```

Okay, so we've deployed Elasticsearch; let's move on to securing it so that we can use the detection engine and Fleet.

## Securing Elasticsearch

As mentioned previously, the Elastic Stack can be looked at as blocks that you assemble in different ways to meet your specific use case. Some people use just Elasticsearch, some use just Logstash, and some use the entire stack. With that in mind, while greatly improved, configuring Elastic is still a bit of a disjointed process. I should say, I don't have a better solution, but that doesn't change the fact that it's still a bit cumbersome.

To get started, let's update the configuration to enable security.

There is a lot in this configuration file, but most of it is either kept as the default or commented out. Most of it is there as a guide (which is helpful). Again, using nano, open the configuration file and add a few lines. Remember, to save and exit in nano, use *Ctrl + X, Y, Enter*:

```
$ sudo nano /etc/elasticsearch/elasticsearch.yml
xpack.security.enabled: true
discovery.type: single-node
network.host: 0.0.0.0
discovery.seed_hosts: ["0.0.0.0"]
xpack.security.authc.api_key.enabled: true
```

After you've made that change, restart Elasticsearch and verify that the service restarts. If not, review the previous steps:

```
$ sudo systemctl restart elasticsearch
$ systemctl status elasticsearch

● elasticsearch.service - Elasticsearch
   Loaded: loaded (/usr/lib/systemd/system/elasticsearch.service; enabled; vendor preset: disabled)
   Active: active (running) since Fri 2021-02-19 04:40:04 UTC;
             7s ago
             ...

```

Next, we need to configure the passphrases for the accounts. As I've mentioned previously, I prefer simplicity for demonstrations and labs, so I make all of these the same simple passphrase. If you have other processes, please feel free to use them.

Elastic provides a utility to allow us to set the passphrases for all of the accounts that we'll be using:

```
$ sudo /usr/share/elasticsearch/bin/elasticsearch-setup-
passwords interactive
```

From here, you'll be asked to set the passphrase for several accounts. **Remember them as we'll need them later.**

Let's test to make sure that we've enabled security and we can authenticate.

Just as before, let's hit the Elasticsearch API with cURL and see what we get:

```
$ curl localhost:9200?pretty
```

We get a 401 error back, which as we can see is an authentication error:

```
...
{
    "type" : "security_exception",
    "reason" : "missing authentication credentials for REST
request [/?pretty]",
    "header" : {
        "WWW-Authenticate" : "Basic realm=\"security\""
charset=\"UTF-8\""
    }
...
...
```

Now, let's try passing the `elastic` username, which we set in the previous step. To pass the username using cURL, we can use the `-u` switch, followed by the username. We should get prompted for a passphrase and then will get the welcome output from Elasticsearch:

```
$ curl -u elastic localhost:9200

Enter host password for user 'elastic':

{
    "name" : "elastic-packetpub.local",
    "cluster_name" : "elasticsearch",
    "cluster_uuid" : "_yiyiDYdQ620Y6FjODu1JQ",
    "version" : {
        "number" : "7.11.1",
        "build_flavor" : "default",
        "build_type" : "rpm",
        "build_hash" : "ff17057114c2199c9c1bbecc727003a907c0db7a",
        "build_date" : "2021-02-15T13:44:09.394032Z",
        "build_snapshot" : false,
        "lucene_version" : "8.7.0",
        "minimum_wire_compatibility_version" : "6.8.0",
        "minimum_index_compatibility_version" : "6.8.0"
    }
}
```

```
"minimum_index_compatibility_version" : "6.0.0-beta1"  
},  
"tagline" : "You Know, for Search"  
}
```

Okay, that's the easy part, an Elasticsearch deployment with basic authentication!

#### Important note

We will not be deploying TLS for Elasticsearch, Kibana, or Beats. TLS is extremely important for production systems to encrypt sensitive traffic between pieces of infrastructure. In a small, non-production lab environment, managing the certificate process falls well beyond the scope. I strongly encourage you to explore the best way for you to deploy TLS in your environment as necessary. The Elastic documentation on configuring and deploying TLS in the Elastic Stack can be found at the following link: <https://www.elastic.co/guide/en/elasticsearch/reference/current/configuring-tls.html>.

In this section, we built a CentOS VM for the Elastic Stack, installed Guest Additions, and deployed and secured Elasticsearch. While this was a long process, having this appropriately done now will make things easier as we progress.

Next, we will install Elastic Agent. This will be used later as the Fleet server, but we want to install it now and configure it later.

## Installing Elastic Agent

On the Elastic VM, Elastic Agent is used to proxy Fleet policies to other enrolled agents. We'll get into Fleet in *Chapter 6, Data Collection with Beats and Elastic Agent*, and Elastic Agent in detail in almost every chapter later in the book.

Still on the command line, simply type the following:

```
$ sudo dnf install elastic-agent  
$ sudo systemctl enable elastic-agent
```

This will install Elastic Agent and configure it to start on boot, but not start it yet. We will configure and start it later in this chapter.

Next, we need to move on to Kibana so that we can access Elasticsearch beyond via the API.

## Installing and configuring Kibana

Now that we've deployed Elasticsearch, we need to build Kibana. A deployment of Kibana is pretty simple, and connecting it to Elasticsearch using basic authentication isn't terribly difficult either.

### Installing Kibana

As we've already installed the Elastic repository, we can simply use that to install Kibana using yum or DNF and enable it to start on boot:

```
$ sudo dnf install kibana  
$ sudo systemctl enable kibana
```

Now that we've installed and configured Kibana to start on boot, we can continue to connect Kibana to Elasticsearch.

### Connecting Kibana to Elasticsearch

Kibana (and Beats for that matter) uses a Java KeyStore to manage and secure credentials. We're going to add `elasticsearch.username` and `elasticsearch.password` to the KeyStore.

This is the username and password used by Kibana to authenticate to Elasticsearch. We set these when we configured all of the credentials during the Elasticsearch setup. `elasticsearch.username` is `kibana_system` and `elasticsearch.password` is something you set:

```
$ sudo /usr/share/kibana/bin/kibana-keystore add elasticsearch.  
username  
Enter value for elasticsearch.username: kibana_system  
  
$ sudo /usr/share/kibana/bin/kibana-keystore add elasticsearch.  
password  
Enter value for elasticsearch.password: *****
```

Next, we need to allow remote connections to Kibana and configure an encryption key for saved objects. To do this, we'll make one change and one addition to the Kibana configuration file.

Like the Elasticsearch configuration file, there is a lot here, but most of it is commented out and left for reference:

- **Update:** For the updated information, we'll uncomment `#server.host: "localhost"` and change `localhost` to `0.0.0.0` to allow external access.
- **Add:** For the addition, we'll add the following field with any random 32 characters. It doesn't matter what those characters are or where you add the field:

```
xpack.encryptedSavedObjects.encryptionKey: "any or more  
32-characters"
```

As we'll be in a non-production development mode and are not using TLS, we need to disable this check for Fleet:

```
xpack.fleet.agents.tlsCheckDisabled: true
```

The Kibana configuration will look like this when completed:

```
# Kibana is served by a back end server. This setting specifies  
the port to use.  
#server.port: 5601  
  
...  
server.host: "0.0.0.0"  
xpack.encryptedSavedObjects.encryptionKey: "thirty-two-or-more-  
random-characters"  
xpack.fleet.agents.tlsCheckDisabled: true  
...
```

Once we've made those two small changes, let's restart Kibana and connect from our browser (finally):

```
$ sudo systemctl restart kibana
```

Now that we've gotten Elasticsearch built and Kibana connected, we should finally be able to use our browser to accelerate into the final configuration steps.

## Connecting to Kibana from a browser

If you'll remember, when we built our Elastic VM, we enabled some port forwarding so that we could connect from our host to our guest. Now is when we'll get to see whether that all worked out.

Still on our Elastic VM, we need to make a few port changes to allow remote access. We can do that using the `firewall-cmd` command.

We'll be adding port 5601 for Kibana, port 9200 for Elasticsearch, and port 8220 for the Fleet server:

```
$ sudo firewall-cmd --add-port=5601/tcp --add-port=9200/tcp  
--add-port=8220/tcp --permanent  
$ sudo firewall-cmd --reload
```

Back on our host machine, open a web browser and browse to `http://localhost:5601`, and you should be presented with a Kibana web interface asking for a username and password:

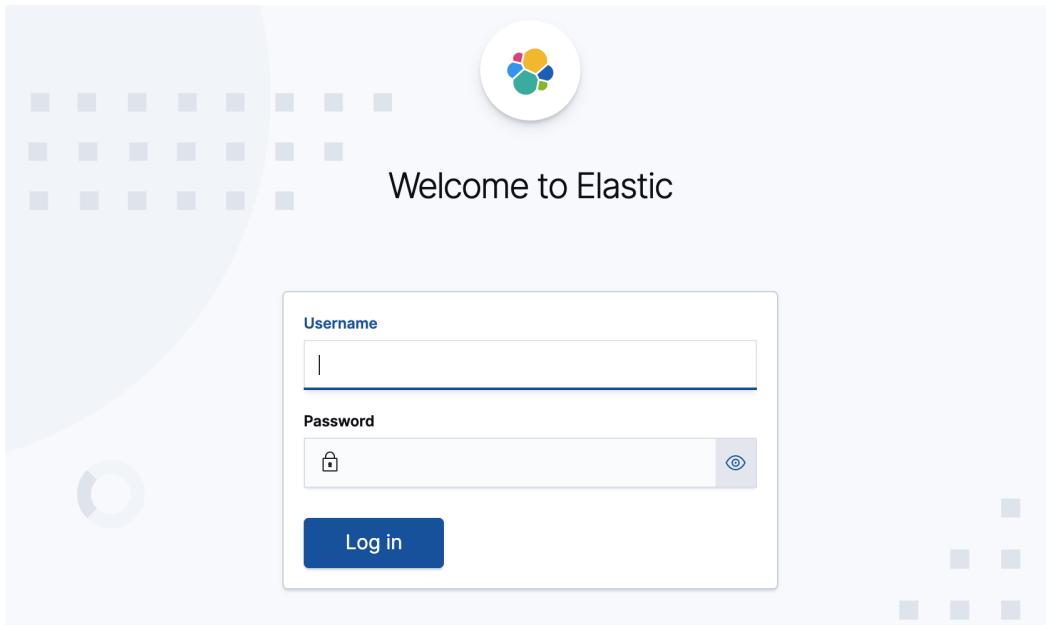


Figure 5.1 – Kibana login page

Let's log in with the `elastic` account we created earlier and prepare to connect our victim machines.

In this section, we installed and configured Kibana to connect to our Elasticsearch node.

Next, we need to enable the detection engine and Fleet so that we can deploy and configure Elastic Agent.

## Enabling the detection engine and Fleet

Now that we've built and configured security for Elasticsearch and Kibana, let's enable the detection engine and Fleet. The detection engine is how we'll ingest and manage the prebuilt Elastic rules for the Security app and Fleet is how we'll centrally manage collection agents.

### Detection engine

The **detection engine** is where prebuilt detection logic is created and managed for the Security app. Detection logic, as utilized in the Security app, is alerts that are generated by certain different conditions on the endpoints. This is not things such as malware alerts, but more like "a binary is being run from the recycle bin." These rules are hand-created by contributors to the Detection Rules GitHub repository (<https://github.com/elastic/detection-rules>). We'll spend more time on Detection Rules in the following chapters.

For now, we want to enable the prebuilt rules:

1. From your browser, log in to Kibana (<http://localhost:5601>) with the Elastic account and passphrase you created in the *Building Elasticsearch* section.
2. Once you're logged in, you can either click on the **Security** tile or click on the hamburger menu and scroll down to **Security Overview**:

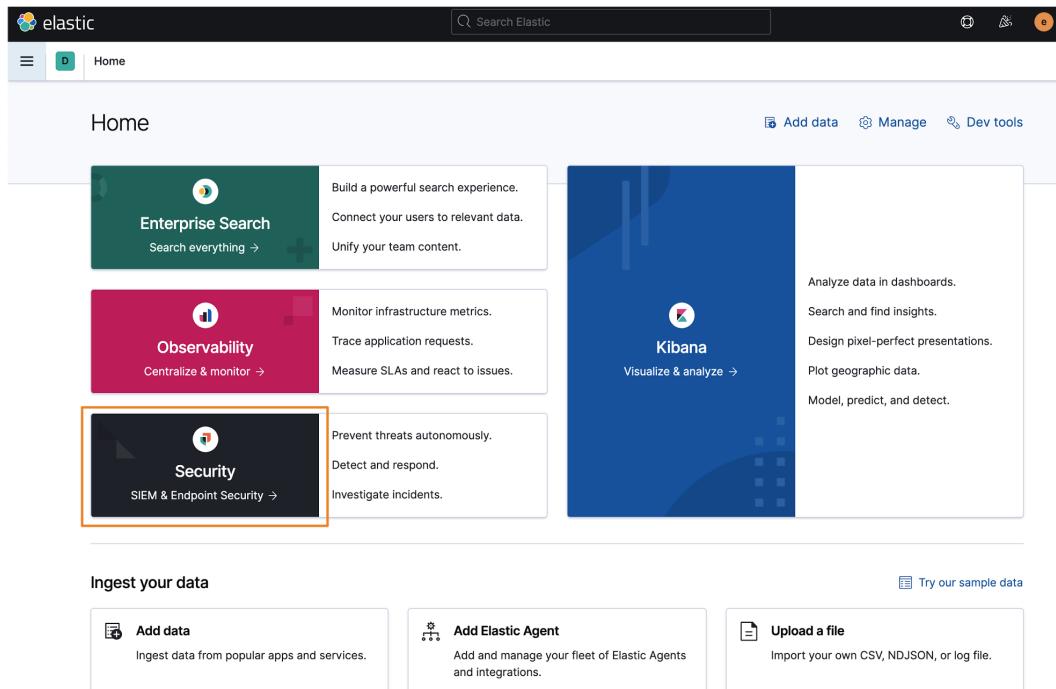


Figure 5.2 – Accessing the Security app

Feel free to explore around here, but for this section, we're just going to be loading the prebuilt rules. We haven't sent any data in yet, so everything will still be blank.

3. Click on the **Detections** tab to open the detection engine:

The screenshot shows the 'Security / Overview' page of the Elastic Security application. The top navigation bar includes the elastic logo, a search bar, and several icons. Below the bar, a secondary navigation menu has tabs for 'Overview' (selected), 'Detections' (highlighted with an orange border), 'Hosts', 'Network', 'Timelines', 'Cases', 'Administration', and a 'Add data' button. The main content area features a large 'Welcome to Elastic Security. Let's get you started.' message with a small icon above it. Below this, there are three callout boxes: one about the Elastic Agent, one about Beats, and one about Endpoint Security, each with a corresponding 'Add data' button.

Welcome to Elastic Security. Let's get you started.

Elastic Security integrates the free and open Elastic SIEM with Endpoint Security to prevent, detect, and respond to threats. To begin, you'll need to add security solution related data to the Elastic Stack. For additional information, you can view our [getting started guide](#).

The Elastic Agent provides a simple, unified way to add monitoring to your hosts.

Add data with Elastic Agent

Lightweight Beats can send data from hundreds or thousands of machines and systems

Add data with Beats

Protect your hosts with threat prevention, detection, and deep security data visibility.

Add Endpoint Security

Figure 5.3 – Security app welcome page

4. Next, click on the blue **Manage detection rules** button on the right:

The screenshot shows the Elastic Security interface for 'Detections'. At the top, there's a search bar labeled 'Search Elastic' and some user icons. Below the header, a navigation bar includes 'Overview', 'Detections' (which is underlined in blue), 'Hosts', 'Network', 'Timelines', 'ML job settings', and a '+ Add data' button. Underneath this is a toolbar with 'Search' (with a dropdown arrow), 'KQL' (with a dropdown arrow), a date range selector ('Last 24 hours'), 'Show dates', and a 'Refresh' button. A 'Filter' icon and a '+ Add filter' link are also present. The main area is titled 'Detection alerts' and contains a 'Trend' section. This section has a 'Stack by' dropdown set to 'signal.rule.name'. Below the dropdown, it says 'Showing: 0 alerts' and 'No data to display'. In the top right corner of the main area, there's a blue button with a gear icon labeled 'Manage detection rules', which is highlighted with an orange rectangle.

Figure 5.4 – The detection engine welcome page

5. Click on **Load Elastic prebuilt rules and timeline templates** to load the prebuilt rules. Again, we'll spend plenty of time in the detection engine rules:

The screenshot shows the 'Detection rules' section of the Elastic Security interface. At the top, there's a navigation bar with tabs for Overview, Detections (which is selected), Hosts, Network, Timelines, Cases, and Administration. Below the navigation bar, there are several buttons: 'Load Elastic prebuilt rules and timeline templates' (highlighted with an orange border), 'Upload value lists', 'Import rule', and 'Create new rule'. Under the 'Detection rules' heading, there are three tabs: Rules (selected), Rule Monitoring, and Exception Lists. A search bar and a 'Tags' dropdown are also present. The main content area is titled 'All rules' and includes a message about prebuilt detection rules. It features two buttons at the bottom: 'Load Elastic prebuilt rules and timeline templates' (highlighted with an orange border) and 'Create your own rules'.

Figure 5.5 – The Detection rules welcome page

This will load hundreds of rules and we'll see a success flyout after a few seconds:

The screenshot shows the Elastic Security app interface. At the top, there's a navigation bar with tabs for Overview, Detections (which is selected), Hosts, Network, Timelines, Cases, and Administration. Below the navigation is a search bar labeled "Search Elastic". The main area is titled "Detection rules" and has tabs for Rules, Rule Monitoring, and Exception Lists. Under the "All rules" heading, a table lists several rules with columns for Rule name, Risk score, Severity, Last run, Last response, Last updated, Version, Tags, and Activated status. One rule is highlighted with a yellow background. A callout box points to the bottom right of the table, containing the text: "✓ Installed pre-packaged rules and timeline templates from elastic".

Figure 5.6 – The detection rules successfully loaded

In this section, we loaded prebuilt Elastic rules into the detection engine. We'll return to this Security app as we continue. Next, we're going to build a Fleet policy that we can deploy into our victim machines.

## Fleet

**Fleet** is the central management hub for deployed collection and protection agents.

### Important note

As of the time of writing, Fleet is still in beta. There could be modifications or changes needed to these steps as Fleet becomes **Generally Available (GA)** in the future. Review Elastic's official Fleet and Elastic Agent documentation for updates: <https://www.elastic.co/guide/en/fleet/current/fleet-overview.html>.

To get to Fleet, click on the hamburger menu and then scroll down to **Fleet**:

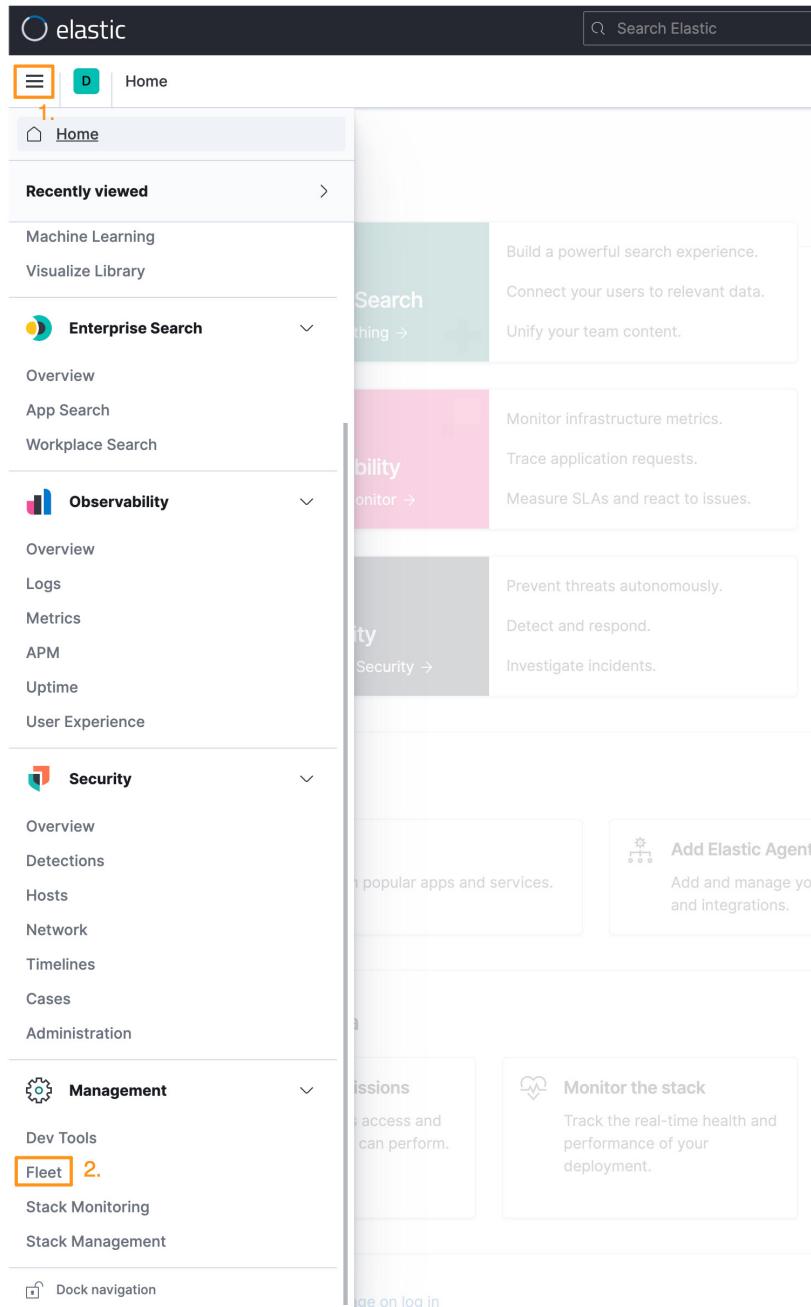


Figure 5.7 – Accessing Fleet

Fleet will take a minute or so to load for the first time. This is a one-time delay.

Once Fleet loads, you'll be on the welcome page. Feel free to explore this app, but for now, we're going to focus on building a Windows and Linux collections and defense policy:

The screenshot shows the Elastic Fleet welcome page. At the top, there's a navigation bar with the Elastic logo, a search bar, and user icons. Below the navigation, a breadcrumb trail shows 'Fleet / Overview'. The main content area has a title 'Fleet BETA' and a subtitle 'Manage Elastic Agents and their policies in a central location.' To the right is a blue button with '+ Add agent'. The page is divided into four main sections:

- Integrations**: Shows 58 available integrations, 2 installed, and 0 updates available.
- Agent policies**: Shows 1 total available policy and 1 used integration.
- Agents**: Shows 0 total agents, 0 active, 0 offline, and 0 error.
- Data streams**: Shows 0 data streams, 0 namespaces, and 0B total size.

Figure 5.8 – Fleet welcome page

Next, in the upper-right corner, you'll see a cog wheel and **Settings**. Click that and change your Fleet Server hosts to `https://172.16.0.103:8220` and Elasticsearch hosts to `http://172.16.0.103:9200` (respectively). Click **Save and apply settings**:

**Note**

Fleet Server and Elasticsearch are using the same IP and that IP is your internet network (`intnet`) IP address. This `intnet` IP was configured by the DHCP steps we completed in *Chapter 5, Building Your Hunting Lab – Part 2*. Additionally, the Fleet Server host is over HTTPS, not HTTP. There is no additional configuration needed to set up HTTPS; it is managed by Fleet. Ensure you use the right IP address as the screenshots here may be using a different IP schema.

## Fleet settings

These settings are applied globally to the `outputs` section of all agent policies and affect all enrolled agents.

**Fleet Server hosts**

`https://172.16.0.3:8220` 

Specify the URLs that your agents will use to connect to a Fleet Server. If multiple URLs exist, Fleet shows the first provided URL for enrollment purposes. Refer to the [Fleet User Guide](#).

**Elasticsearch hosts**

`http://172.16.0.3:9200` 

Specify the Elasticsearch URLs where agents send data.

**Elasticsearch output configuration (YAML)**

`# YAML settings here will be added to the Elasticsearch output section of each pol`

Figure 5.9 – Fleet settings

This setting will configure Elastic Agent so that it is reporting to Kibana for configuration updates as well as sending its data into the proper Elasticsearch instance.

Next, we need to collect an enrollment token. This enrollment token will be used to configure an Elastic Agent as an actual Fleet server. This Elastic Agent will run on the same Elastic server and handle the management of other Elastic Agents. All data will still be sent directly to Elasticsearch.

Click on the **Agents** tab and then click on **Generate service token**:

The screenshot shows the Elastic Fleet interface with the 'Agents' tab selected. A modal window titled 'Add a Fleet Server' is open. The modal contains three steps: 1. Download the Elastic Agent to your host, with a link to the download page. 2. Generate a service token, with a prominent blue button labeled 'Generate service token' highlighted by an orange box. 3. Start Fleet Server. Below the steps, a status message says 'Waiting for a Fleet Server to connect...'.

Figure 5.10 – Add a Fleet Server

Next, you'll be presented with a token. You shouldn't need this again for our lab, but you can copy it down.

Under the service token, you'll be asked what platform you want to configure Elastic Agent for. We're using CentOS and installed Elastic Agent using DNF in a previous step, so select **RPM / DEB** from the dropdown and copy the enrollment syntax:

## Add a Fleet Server

A Fleet Server is required before you can enroll agents with Fleet. See the [Fleet User Guide](#) for more information.

- 1 Download the Elastic Agent to your host**

You can download the agent binaries and their verification signatures from the Elastic Agent download page.

[Go to download page](#)
- 2 Generate a service token**

A service token grants Fleet Server permissions to write to Elasticsearch.

Save your service token information. This will be shown only once.

**Service token** AAEAAWVsYXN0aWVmZmx1ZXQtc2VydmVyL3Rva2VuLTE2MjIwOTE5NzE1ODU6T31ZdGh4M1d [Copy](#)
- 3 Start Fleet Server**

From the agent directory, copy and run the appropriate quick start command to start an Elastic Agent as a Fleet Server using the generated token and a self-signed certificate. See the [Fleet User Guide](#) for instructions on using your own certificates for production deployment. All commands require administrator privileges.

Platform [RPM / DEB](#) [▼](#)

```
sudo elastic-agent enroll -f --fleet-server-es=http://172.16.0.3:9200 --1 Copy
```

If you are having trouble connecting, see our [troubleshooting guide](#).

 Waiting for a Fleet Server to connect...

Figure 5.11 – Fleet enrollment command

Next, we'll leave Kibana and complete the enrollment on the command line of the Elastic VM.

## Enrolling Fleet Server

On the command line of the Elastic VM, we simply need to run the command that we copied in the previous step (it will include your assigned service token). On the command line, type the following (your IP address may be different):

```
$ sudo elastic-agent enroll -f --fleet-server-
es=http://172.16.0.3:9200 --fleet-server-service-token=your-
service-token
```

Finally, let's restart Elastic Agent as a Fleet server and show it enrolling into Kibana. Still on the command line, type the following:

```
sudo systemctl restart elastic-agent
```

Next, back in Kibana, on the Fleet **Agents** page, we should see Elastic Agent enroll as a Fleet server. The agent could take a few minutes to check in:

The screenshot shows the Kibana interface with the 'elastic' dashboard selected. The top navigation bar includes 'Fleet' and 'Agents'. The main content area is titled 'Agents' and displays a table of agent status. The table has columns for Host, Status, Agent policy, Version, Last activity, and Actions. One row is visible for 'elastic-packetpub.local', which is 'Healthy' (green), using 'Default Fleet Server policy' (rev. 4), version 7.13.0, and last active 31 seconds ago. A blue button labeled '+ Add agent' is located in the top right corner of the table area.

Host	Status	Agent policy	Version	Last activity	Actions
elastic-packetpub.local	Healthy	Default Fleet Server policy rev. 4	7.13.0	31 seconds ago	...

Figure 5.12 – Elastic Agent enrolled as a Fleet server

In this section, we enabled and configured both the detection engine and Fleet. Both of these Kibana features will be paramount when we get into the Elastic Security app in *Chapter 8, The Elastic Security App*.

Now that we have built our storage and analysis platform (the Elastic Stack), we need to build the victim machine that we'll use to collect data from.

# Building a victim machine

In this section, we'll be building a machine that will be used to collect data from. While collecting normal system information is valuable, we'll be collecting security-relevant data from these systems. We don't want to detonate malware or perform risky behavior on a production system, so we'll be making a system purely to generate malicious data for us to analyze. We also call these victim machines.

In this section, we'll build one victim machine. Feel free to mix and match this approach with more than one Windows or Linux machine, use a different version of Windows or Linux, or if you're running low on resources, pick one or the other instead of both.

## Collecting the operating systems

First, we need to collect the operating system ISO images for Windows.

### Windows

Microsoft uses the Evaluation Center to provide 90-day copies of their software to IT professionals for zero cost. These are not meant for production deployment as their functionality will be reduced after the 90 days. For a long-term strategy, you should consider purchasing a license for Windows.

Browse to the Evaluation Center (<https://www.microsoft.com/en-us/evalcenter/evaluate-windows-10-enterprise>) and download the **ISO - Enterprise** version of Windows.

Now that we've collected the operating system, let's build it into a VM in VirtualBox.

## Creating the virtual machine

In these next steps, we'll install the operating system, perform updates, and configure the guest additions.

To get started with Windows, let's open VirtualBox and click on the **New** icon. Use the same steps that you used when configuring the Elastic VM previously. Note that you'll not need to forward any ports for the Windows VM.

Of note, for Windows, we'll be using Network Adapter 1 and Network Adapter 3:

- **Name:** Windows 10 Victim Box (this can be anything you want).
- **Machine Folder:** This should be pre-populated, but you can adjust it if needed.

- **Type:** Microsoft Windows.
- **Version:** Windows 10 (64-bit).
- 4,192 MB RAM.
- 30 GB hard disk (feel free to increase this if you have the resources).
- Set the boot order to **Hard Disk** then **Optical**.
- **Network:** Adapter 1 – Internet.

Remember to attach your Windows ISO:

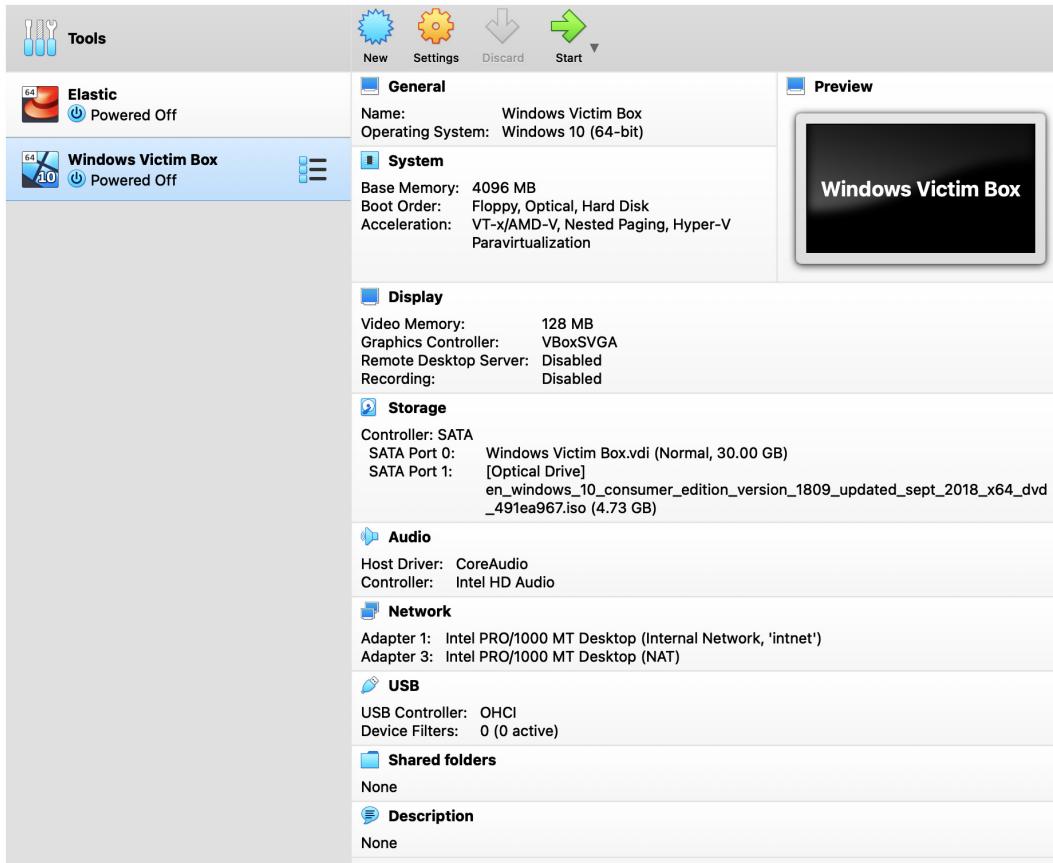


Figure 5.13 – Windows virtual machine details

Now that we have built the VM, we can proceed with installing the operating system.

## Installing Windows

First, we'll install Windows. The Windows VM is largely using default settings. There is a bit of a cumbersome process of setting up a local account instead of using an online Microsoft account, but there's a pretty simple workaround we'll use.

In VirtualBox, select the Windows VM and click the green **Start** button. Your VM should start up into the Windows installation process.

You should select the default options until you get to the section to enter your product key. As we're using an evaluation version, we can simply select **I don't have a product key** to proceed:

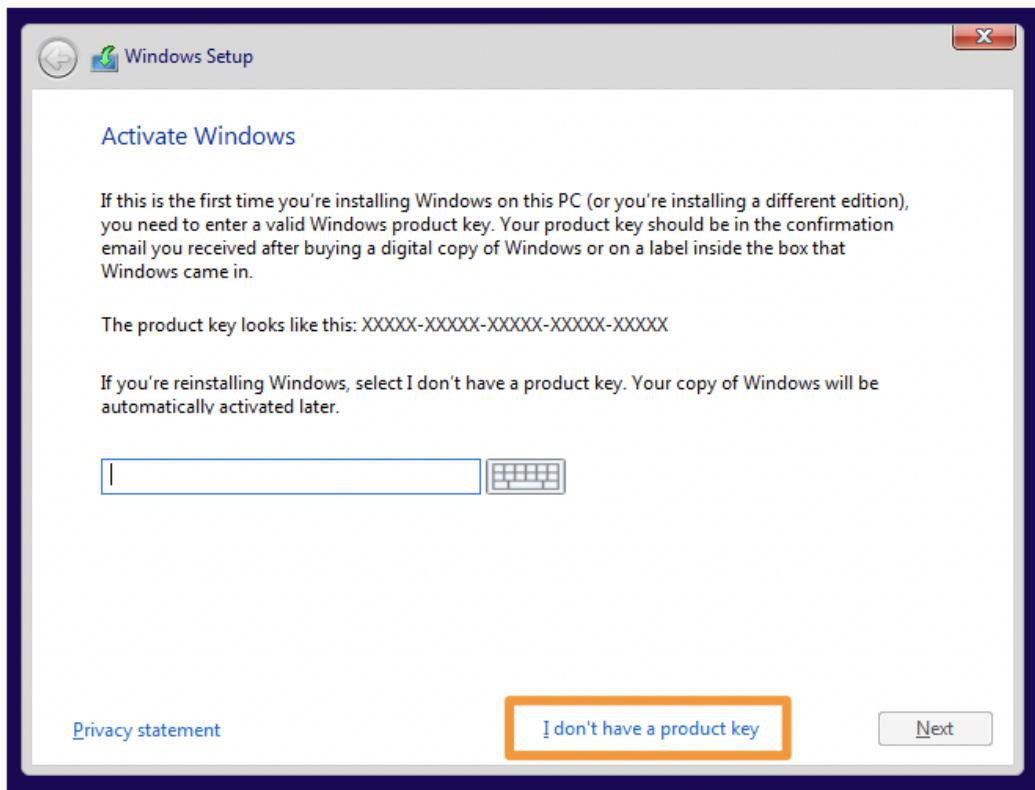


Figure 5.14 – Windows virtual machine product key selection

Continue with your default selections until you are asked what type of installation you need. Here, select **Custom: Install Windows only (advanced)**:

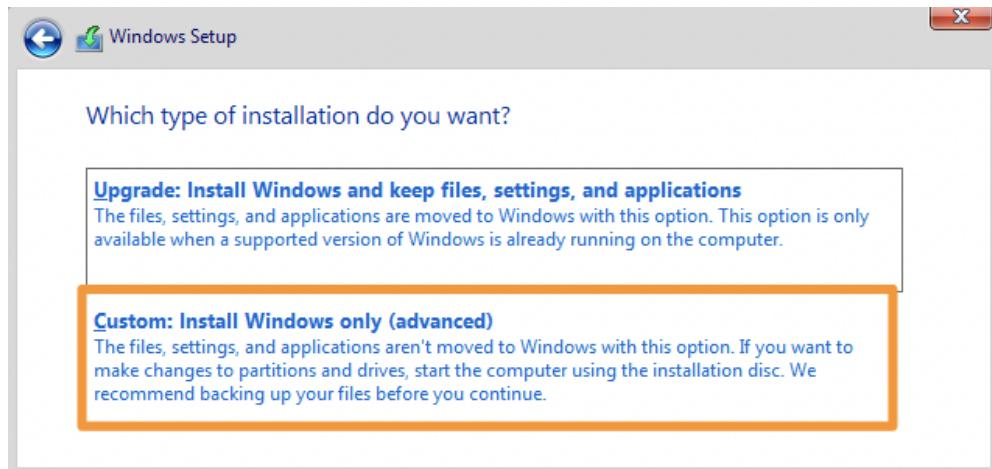


Figure 5.15 – Windows virtual machine installation type

Allow the Windows installation to continue; it will take a few minutes. Afterward, the VM will reboot a few times.

Eventually, you'll get to select your region and keyboard layout. Select whatever keyboard layout you prefer.

When you get to the network configuration, click in the lower left on **Skip for now**. We don't have a network connection to the internet yet. You'll get asked a second time to configure the network; select **No**:

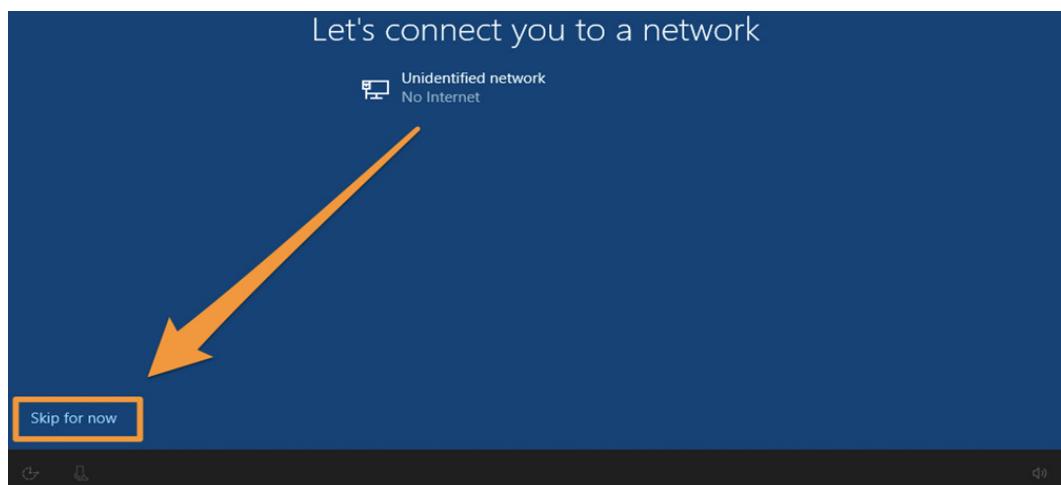


Figure 5.16 – Skipping the network connection

Now that you have bypassed using the internet to set up an account, you can move on to create a local account:

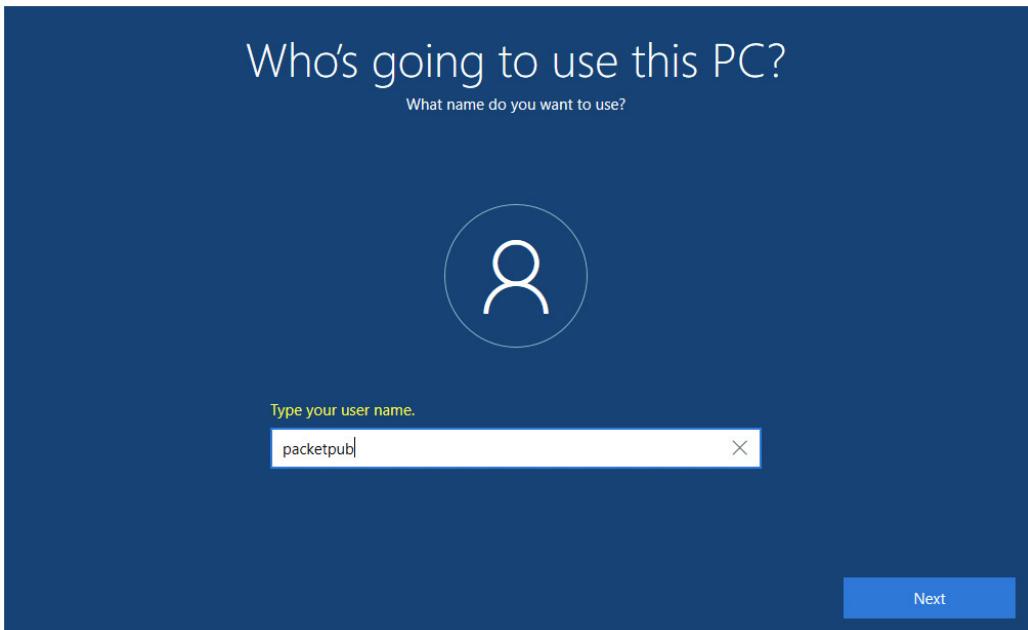


Figure 5.17 – Windows virtual machine – create your account

Create your account name, passphrase, and security questions. Decline enabling Cortana, decline the activity history, and uncheck all of the privacy settings. Finally, click on **Accept** and allow the installation to complete.

Once you're completed the installation, using the same steps highlighted previously, re-enable Network Adapter 3.

Finally, let's do a test to check to make sure that the Windows guest can reach the Elastic box.

## Connection test

First, turn on your Elastic VM (if it isn't already running). Do a local test, from the Elastic VM, to make sure Elasticsearch is running properly (`curl -u elastic localhost : 9200`). If you don't get the Elasticsearch welcome message, go back to the *Building Elasticsearch* section and ensure you've followed all the steps.

Next, from the Windows VM, open a terminal window (`cmd.exe` or `powershell.exe`) and check to see whether you can connect with `curl -u elastic 172.16.0.3:9200` (remember, to use your intnet IP address as it may differ from the example). You should get the Elasticsearch welcome message!

```
Microsoft Windows [Version 10.0.17763.107]
(c) 2018 Microsoft Corporation. All rights reserved.

C:\Users\packetpub>curl -u elastic 192.168.1.101:9200
Enter host password for user 'elastic':
{
  "name" : "elastic-packetpub.local",
  "cluster_name" : "elasticsearch",
  "cluster_uuid" : "_iyidYdQ620Y6FjODU1JQ",
  "version" : {
    "number" : "7.11.1",
    "build_flavor" : "default",
    "build_type" : "rpm",
    "build_hash" : "ff17057114c2199c9c1bbecc727003a907c0db7a",
    "build_date" : "2021-02-15T13:44:09.394032Z",
    "build_snapshot" : false,
    "lucene_version" : "8.7.0",
    "minimum_wire_compatibility_version" : "6.8.0",
    "minimum_index_compatibility_version" : "6.0.0-beta1"
  },
  "tagline" : "You Know, for Search"
}

C:\Users\packetpub>
```

Figure 5.18 – Windows virtual machine connecting to Elasticsearch

Now that we have a fully functional Windows VM, we should install Guest Additions to make the VM experience a bit smoother.

## Installing VirtualBox Guest Additions

Unlike when installing Guest Additions on Linux, we don't need to hop around with adding and removing virtual disks. We can simply select **Devices** in the VM menu, and then **Insert Guest Additions CD image....**

Back in the Windows VM, open the D :\ drive, right-click on **VBox Windows Additions**, and select **Run as administrator**. Select all of the defaults and then shut down the machine.

## Enabling the clipboard

While the Windows VM is still shut down, click on **Settings** and then **Advanced**, and then set **Shared Clipboard** to **Host To Guest** so we can copy data into the VM:

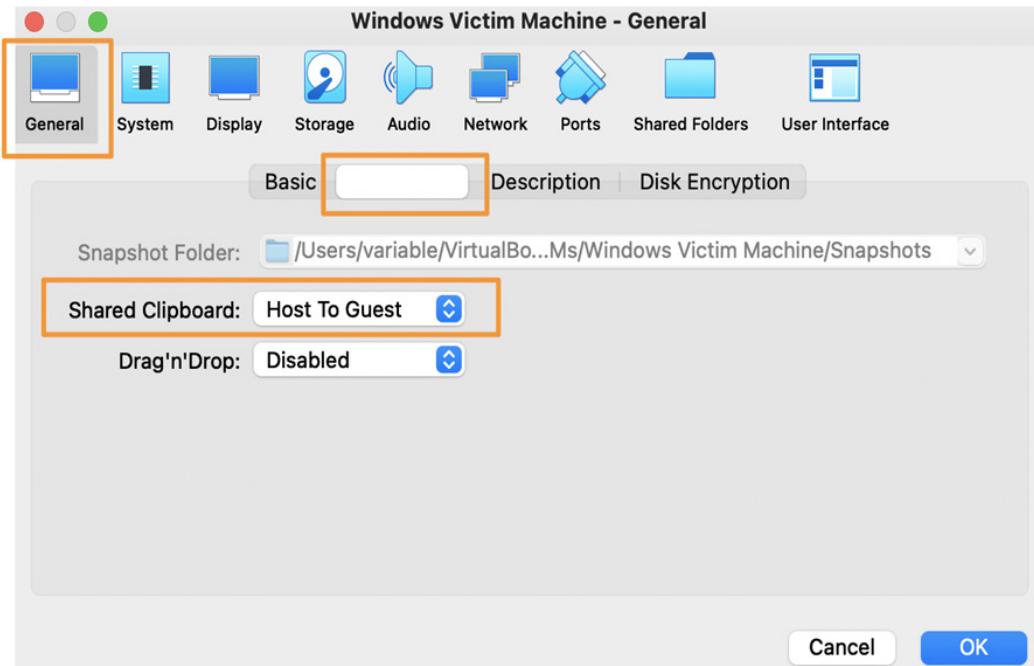


Figure 5.19 – Enabling Shared Clipboard

In the preceding screenshot, we're enabling the ability to copy data into the VM using the clipboard.

## Enabling Network Adapter 3

Now we're going to add an additional network adapter. This will be Adapter 3. At the end of the configuration, you'll have two network adapters enabled: Adapter 1 and Adapter 3.

While your Windows VM is shut down, click on **Settings** and then **Network** and enable Adapter 3. Ensure it is set to **NAT**:

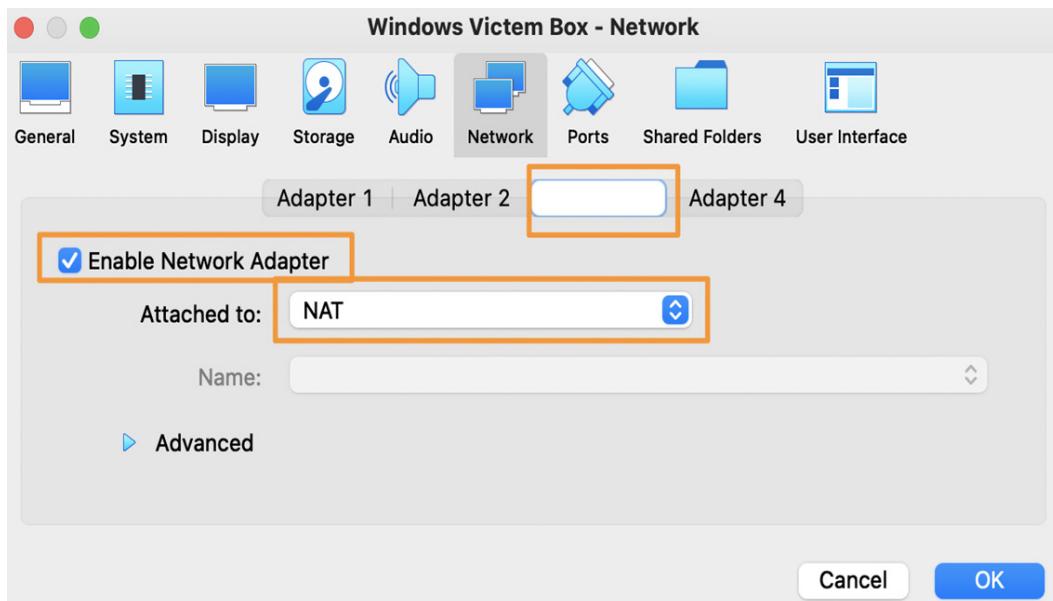


Figure 5.20 – Enabling Adapter 3

In the preceding screenshot, we're enabling Adapter 3 on the Windows VM.

## Windows Update

I recommend letting your Windows updates run. This isn't necessary, but I've found that letting this process complete makes the process a bit easier versus trying to battle for resources.

If you choose not to run updates, that's totally fine, but I do recommend downloading the **new Edge** (or Chrome, Firefox, or Safari) browser because the default Edge has a hard time accessing some of the pages that we'll need to download additional packages.

## Last mile configurations

Finally, we'll be doing some "last mile" configurations to ensure that we're collecting the most amount of value from the PowerShell logs.

**PowerShell script block logging** records the content of all the script blocks that it processes. This is very valuable in tracking the malware and adversaries that leverage PowerShell.

Enabling `ScriptBlockLogging` is accomplished with a simple PowerShell function:

1. Open PowerShell as an administrator.
2. Paste the following PowerShell function into the PowerShell window:

```
function Enable-PSScriptBlockLogging
{
    $basePath = 'HKLM:\Software\Policies\Microsoft\
Windows' +
    '\PowerShell\ScriptBlockLogging'

    if(-not (Test-Path $basePath))
    {
        $null = New-Item $basePath -Force
    }

    Set-ItemProperty $basePath -Name
    EnableScriptBlockLogging -Value "1"
}
```

The preceding code will enable PowerShell script block logging, which will allow us to record detailed information about PowerShell activities on the victim machine.

Now that we have built our victim machine, let's learn how we can use Filebeat to import threat data.

## Filebeat Threat Intel module

**Filebeat** has a Threat Intel module that is intended to import threat data from various feeds. We'll set up three of the feeds that do not require any third-party accounts, but you can set those up as well if you have accounts.

In Elastic 7.12, the Threat Intel module collects data from five sources:

- Abuse Malware
- Abuse URL
- Anomali Limo
- AlienVault OTX (free account required)
- MISP (additional infrastructure required)

We'll go through the steps to set up Abuse Malware, Abuse URL, and Anomali Limo:

1. Log in to the command line of your Elastic VM and install filebeat using DNF:

```
sudo dnf install filebeat -y
```

Once you have installed Filebeat, you need to update the configuration and start collecting data.

2. First, let's enable the Threat Intel Filebeat module. We can simply run the following:

```
sudo filebeat modules enable threatintel
```

3. Now that we've enabled the Threat Intel module, we just need to do a few tweaks in the module configuration:

```
sudo nano /etc/filebeat/modules.d/threatintel.yml
```

4. Once we're in the configuration file, change the settings for MISP and OTX from true to false. Under the Anomali section, uncomment the username and password:

```
# Module: threatintel
# Docs: https://www.elastic.co/guide/en/beats/
filebeat/7.x/filebeat-module-threatintel.html

- module: threatintel
  abuseurl:
    enabled: true

  # Input used for ingesting threat intel data.
  var.input: httpjson

  # The URL used for Threat Intel API calls.
  var.url: https://urlhaus-api.abuse.ch/v1/urls/recent/

  # The interval to poll the API for updates.
  var.interval: 10m

  abusemalware:
    enabled: true
```

```
# Input used for ingesting threat intel data.
var.input: httpjson

# The URL used for Threat Intel API calls.
var.url: https://urlhaus-api.abuse.ch/v1/payloads/
recent/

# The interval to poll the API for updates.
var.interval: 10m

misp:
  enabled: false

# Input used for ingesting threat intel data,
# defaults to JSON.
var.input: httpjson

# The URL of the MISP instance, should end with "/events/restSearch".
var.url: https://SERVER/events/restSearch

# The authentication token used to contact the MISP API. Found when looking at user account in the MISP UI.
var.api_token: API_KEY

# Configures the type of SSL verification done, if MISP is running on self signed certificates
# then the certificate would either need to be trusted, or verification_mode set to none.
#var.ssl.verification_mode: none

# Optional filters that can be applied to the API for filtering out results. This should support the majority of fields in a MISP context.
# For examples please reference the filebeat module documentation.
#var.filters:
```

```
# - threat_level: [4, 5]
# - to_ids: true

# How far back to look once the beat starts up
# for the first time, the value has to be in hours. Each
# request afterwards will filter on any event newer
# than the last event that was already ingested.
var.first_interval: 300h

# The interval to poll the API for updates.
var.interval: 5m

otx:
    enabled: false

# Input used for ingesting threat intel data
var.input: httpjson

# The URL used for OTX Threat Intel API calls.
var.url: https://otx.alienvault.com/api/v1/
indicators/export

# The authentication token used to contact the OTX
# API, can be found on the OTX UI.
var.api_token: API_KEY

# Optional filters that can be applied to retrieve
# only specific indicators.
#var.types: "domain,IPv4,hostname,url,FileHash-
#SHA256"

# The timeout of the HTTP client connecting to the
# OTX API
#var.http_client_timeout: 120s

# How many hours to look back for each request,
# should be close to the configured interval. Deduplication
# of events is handled by the module.
```

```
var.lookback_range: 1h

# How far back to look once the beat starts up for
the first time, the value has to be in hours.

var.first_interval: 400h

# The interval to poll the API for updates
var.interval: 5m

anomali:
  enabled: true

# Input used for ingesting threat intel data
var.input: httpjson

# The URL used for Threat Intel API calls. Limo has
multiple different possibilities for URL's depending
# on the type of threat intel source that is needed.
var.url: https://limo.anomali.com/api/v1/taxii2/
feeds/collections/313/objects

# The Username used by anomali Limo, defaults to
guest.
var.username: guest

# The password used by anomali Limo, defaults to
guest.
var.password: guest

# How far back to look once the beat starts up for
the first time, the value has to be in hours.

var.first_interval: 400h

# The interval to poll the API for updates
var.interval: 5m
```

5. Now that we've configured our module, we need to configure Filebeat itself to write to Elasticsearch. If you remember, we set a passphrase for Elasticsearch, so let's update it in the Filebeat configuration:

```
sudo nano /etc/filebeat/filebeat.yml
```

6. Go down to the Elasticsearch Output section, uncomment out the username and password, and update your password. Save and exit this file:

```
# --- Elasticsearch Output ---
output.elasticsearch:
  # Array of hosts to connect to.
  hosts: ["localhost:9200"]

  # Protocol - either 'http' (default) or 'https'.
  #protocol: "https"

  # Authentication credentials - either API key or
  #username/password.
  #api_key: "id:api_key"
  #username: "elastic"
  #password: "password"
```

These configuration files are available in the chapter's *Technical requirements* section.

7. Now that we've updated this configuration file, let's push our dashboards and ingest pipelines into Elasticsearch:

```
sudo filebeat setup
```

8. Finally, we can start Filebeat and set it to start on boot:

```
sudo systemctl enable filebeat
sudo systemctl start filebeat
```

Let's go into Kibana and check to see whether our data is flowing into the Discover app.

Log in to Kibana, click on **Discover App** on the left, and change your index pattern to `filebeat-*`, and you'll see your Threat Intel data flooding in! We'll use this in *Chapter 8, The Elastic Security App*.

In this section, we finalized our configuration of Elasticsearch, Kibana, and the guests. This will set us up nicely for future chapters.

## Summary

I agree, there was a lot covered in these last two chapters. There are a lot of moving pieces to get a working lab, between the victim machine and the collection and analysis platform. I thought about doing some "figurative hand waving" with a lot of what I have seen in these kinds of kits: "Step 1: Install Elasticsearch; Step 2: Install Windows; Step 3: Profit." I observe those kinds of guides frequently and find that that approach misses a lot of crucial details; so, we went from the ground up, through every step. While that may seem slow for some that are experienced, it's important to get this right or the rest of the book isn't going to be a lot of fun if all you can do is read about how this *could* work. Hands-on for the win!

In the next chapter, we will configure the systems that will collect our data from the victim machine and store it in Elasticsearch.

## Questions

As we conclude, here is a list of questions for you to test your knowledge regarding this chapter's material. You will find the answers in the *Assessments* section of the *Appendix*:

1. What port does Kibana run on?
  - a. 9200
  - b. 5601
  - c. 8220
  - d. 8080
2. What port does Elasticsearch run on?
  - a. 5601
  - b. 22
  - c. 443
  - d. 9200
3. What does Kibana use to store Elasticsearch authentication credentials?
  - a. KeyStore
  - b. Shadow file
  - c. Vault
  - d. Password file

4. What is the name of the agent central management tool in Kibana?
  - a. Detection engine
  - b. Uptime
  - c. Fleet
  - d. Heartbeat
5. What does script block logging do?
  - a. Records the content of PowerShell script blocks
  - b. Logs endpoint network traffic
  - c. Monitors access to the Elasticsearch API
  - d. Logs roles and users in Kibana

## Further reading

To learn more about the subject, check out the following links:

- About logging in Windows: [https://docs.microsoft.com/en-us/powershell/module/microsoft.powershell.core/about/about\\_logging\\_windows?view=powershell-7.1](https://docs.microsoft.com/en-us/powershell/module/microsoft.powershell.core/about/about_logging_windows?view=powershell-7.1)
- Elasticsearch reference: <https://www.elastic.co/guide/en/elasticsearch/reference/current/index.html>
- Kibana reference: <https://www.elastic.co/guide/en/kibana/current/index.html>
- Fleet reference: <https://www.elastic.co/guide/en/kibana/current/fleet.html>

# 6

# Data Collection with Beats and Elastic Agent

In the last chapter, we built the **virtual machines (VMs)** needed for your hunting lab. In this chapter, we're going to configure all of the infrastructure used to collect all of the data we're going to generate once we start threat hunting.

It is important that the two VMs you built in *Chapter 4, Building Your Hunting Lab – Part 1* and *Chapter 5, Building Your Hunting Lab – Part 2*, are operational and are able to communicate using the connection test at the end of the chapter.

In this chapter, you'll learn how to configure the collection agents and tools you installed in *Chapter 4, Building Your Hunting Lab – Part 1* and *Chapter 5, Building Your Hunting Lab – Part 2*. Additionally, we'll cover the configuration of Fleet that we'll use to manage Elastic Agent.

In this chapter, we'll go through the following topics:

- Data flow
- Configuring Winlogbeat and Packetbeat
- Configuring Sysmon for endpoint collection
- Configuring Elastic Agent collection policies and integrations
- Deploying Elastic Agent for collection

## Technical requirements

In this chapter, you will need to have access to the following:

- The Elastic and Windows VMs built in *Chapter 4, Building Your Hunting Lab – Part 1*
- A modern web browser with a UI

The code for the examples in this chapter can be found at the following GitHub link:

[https://github.com/PacktPublishing/Threat-Hunting-with-Elastic-Stack/tree/main/chapter\\_6\\_data\\_collection\\_with\\_beats\\_and\\_the\\_elastic\\_agent](https://github.com/PacktPublishing/Threat-Hunting-with-Elastic-Stack/tree/main/chapter_6_data_collection_with_beats_and_the_elastic_agent).

Check out the following video to see the Code in Action:

<https://bit.ly/3hKyTs4>

## Data flow

Before we get started with data collection, it would be good to have a basic visualization to highlight the data flow for the Elastic Endpoint Agent, Beats, Elasticsearch, Kibana, and Fleet.

In the following diagram, you can see these flows:

- The Elastic Endpoint Agent sends logs to Elasticsearch.
- The Beats (we're using Winlogbeat and Packetbeat, but all Beats do this by default) send their logs to Elasticsearch.
- Elasticsearch data is rendered by Kibana.
- Kibana uses Fleet to send command-and-control instructions to the Elastic Endpoint Agent:

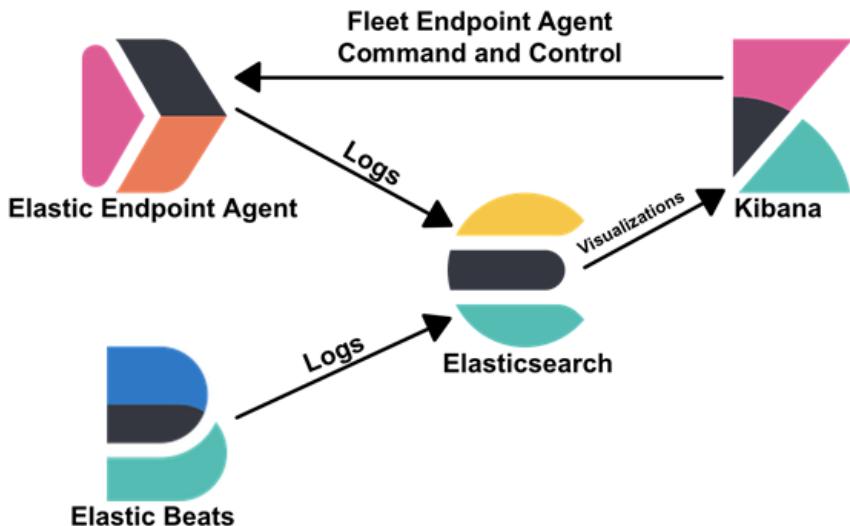


Figure 6.1 – Example of log and Fleet data flow

In this section, we had a high-level exploration of the data flow that we'll use throughout this chapter and the rest of the book.

In the next section, we'll begin configuring our Beats, Winlogbeat and Packetbeat.

## Configuring Winlogbeat and Packetbeat

As stated in *Chapter 3, Introduction to the Elastic Stack*, **Winlogbeat** is a data shipper for Windows events and **Packetbeat** is a data shipper for application-type network events. Both provide a tremendous amount of information by tracking endpoint and network data when threat hunting.

### Installing Beats

Let's download Winlogbeat and Packetbeat, apply some configuration, and run them as services. First, we need to collect the binaries, so from the Windows VM, do the following:

- Download Winlogbeat (the x64 ZIP file, not the MSI version): <https://www.elastic.co/downloads/beats/winlogbeat>.
- Download Packetbeat (the x64 ZIP file, not the MSI version): <https://www.elastic.co/downloads/beats/packetbeat>.
- Download Npcap (**Npcap**, not Nmap): <https://nmap.org/npcap/>.

Find the `winlogbeat-{version}-windows-x86_64.zip` file that was downloaded, right-click it, and extract it to the `c:\Program Files` directory.

Next, we need to create a Java KeyStore so our Beat can authenticate to Elasticsearch. Just like when we did this for Kibana, you'll need to use the `elastic` username and password you set during the *Securing Elasticsearch* section of *Chapter 5, Building Your Hunting Lab – Part 2*.

Open up a terminal window and create the KeyStore. You need administrator privileges, so right-click and select **Run as administrator** when you open Command Prompt (`cmd.exe`). Browse into the Winlogbeat directory (for example, `cd c:\Program Files\Winlogbeat-{version}-windows-x86_64`) and then type the following:

```
winlogbeat.exe keystore create
```

Now that we've created the KeyStore, let's add the credentials to it. This should be the password for the `elastic` user. We're going to call this credential `ES_PWD`:

```
winlogbeat.exe keystore add ES_PWD
```

```
Enter value for ES_PWD: [password for the elastic user]
```

```
Successfully updated the keystore
```

Next, let's open up the Winlogbeat configuration file and tell it where to send its data and what credentials to use.

Using `Notepad.exe` (or similar), open up `elasticsearch.yml`. This file is located in the current Winlogbeat directory (`c:\Program Files\winlogbeat-{version}-windows-x86_64`). You may want to open this from your administrative Command Prompt as this file requires a privileged account to modify (`notepad.exe winlogbeat.yml`).

Scroll down to the Kibana section. Uncomment the `host` line and change the IP address to `172.16.0.3`:

```
# === Kibana ===

# Starting with Beats version 6.0.0, the dashboards are loaded
via the Kibana API.

# This requires a Kibana endpoint configuration.

setup.kibana:

# Kibana Host
```

```
# Scheme and port can be left out and will be set to the
# default (http and 5601)
# In case you specify an additional path, the scheme is
# required: http://localhost:5601/path
# IPv6 addresses should always be defined as: https://
[2001:db8::1]:5601
host: "172.16.0.3:5601"172.16.0.3
...
```

Scroll down to the Outputs section. We're going to update the location of Elasticsearch and the credentials. Elasticsearch's IP address is 172.16.0.3, still over port 9200. The username is elastic and the password will be the KeyStore variable, \${ES\_PWD}. Ensure you uncomment username and password.

The configuration will look like this:

```
# === Outputs ===
# Configure what output to use when sending the data collected
# by the beat.
# --- Elasticsearch Output ---
output.elasticsearch:
    # Array of hosts to connect to.
    hosts: ["172.16.0.3:9200"]
    # Protocol - either 'http' (default) or 'https'.
    #protocol: "https"
    # Authentication credentials - either API key or username/
    #password.
    #api_key: "id:api_key"
    username: "elastic"
    password: "${ES_PWD}"
```

Once you've made that change, let's test the configuration and connection to Elasticsearch (ensure your Elastic VM is running):

```
winlogbeat.exe test config
Config OK

winlogbeat test output
elasticsearch: http://172.16.0.3:9200...
    parse url... OK
```

```
connection...
parse host... OK
dns lookup... OK
addresses: 172.16.0.3
dial up... OK
TLS... WARN secure connection disabled
talk to server... OK
version: 7.11.1
```

Looks good, so let's try to manually start Winlogbeat to set up the index patterns and load the dashboards:

```
winlogbeat.exe setup
```

If that runs without error (which it should), let's send in some data:

```
winlogbeat.exe
```

Leave that running, and let's go over to Kibana and see if there is any data. Click on the **Discover** tab, and you should have data pouring in! Remember to make sure your index pattern is set to **winlogbeat-\***:

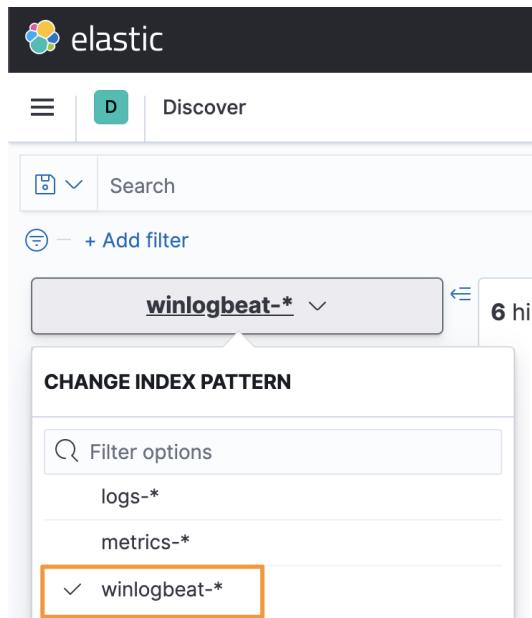


Figure 6.2 – Setting the index pattern to Winlogbeat

Now that we've got Winlogbeat properly configured, we should see the Windows data populating Kibana:

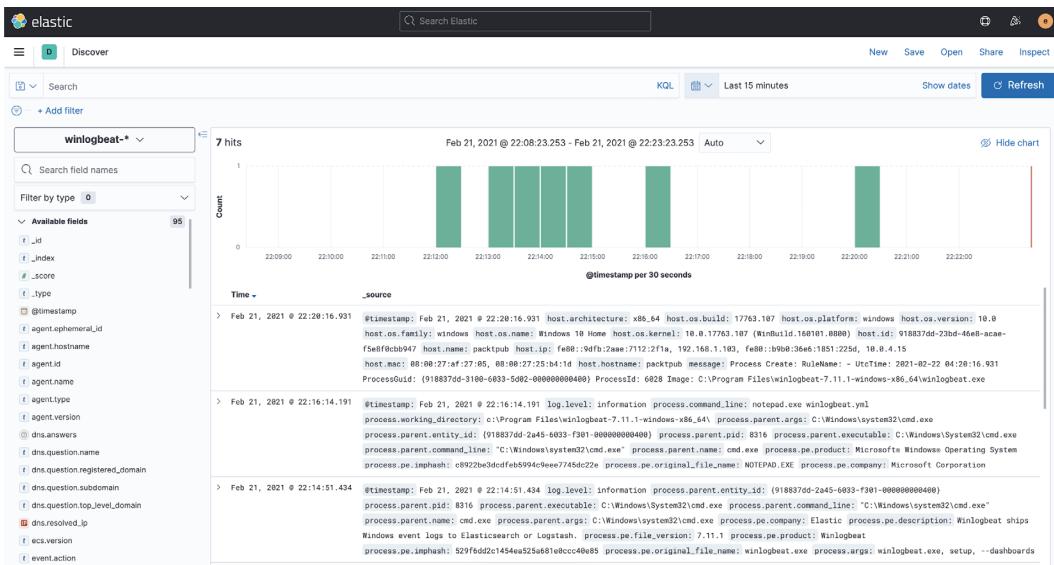


Figure 6.3 – Viewing Windows logs in Kibana

Now that we've proven it's working, let's stop the binary and install it as a service:

```
PowerShell.exe -ExecutionPolicy UnRestricted -File .\install-service-winlogbeat.ps1
```

There is one final step – we need to move the KeyStore that we created into the hidden ProgramData directory. In an administrative terminal, copy the file there:

```
copy "C:\Program Files\winlogbeat-{version}-windows-x86_64\data\winlogbeat.keystore" c:\ProgramData\winlogbeat\
```

#### Important note

C:\ProgramData is a hidden folder by default. You can unhide this in Windows or simply use the preceding command to copy the KeyStore directly into the folder (recommended).

This will install Winlogbeat as a service. It will start on reboot, and we'll be doing that in a bit.

You'll want to repeat all of these steps for Packetbeat, but note that there are a few additional steps.

First, install Npcap (which you downloaded at the beginning of this section). Make sure you got Npcap, not Nmap. You'll need this to load the drivers needed to capture network traffic in Windows.

Next, extract the Packetbeat archive using the same steps you did for Winlogbeat.

Don't forget to create your KeyStore from an administrative terminal:

```
packetbeat.exe keystore create
```

Now that we've created the KeyStore, let's add the credentials to it. This should be the password for the `elastic` user. We're going to call this credential `ES_PWD`:

```
packetbeat.exe keystore add ES_PWD
```

```
Enter value for ES_PWD: [password for the elastic user]
```

```
Successfully updated the keystore
```

You'll want to define the device you want to monitor with Packetbeat. To do this, with an administrative terminal window, run the following:

```
packetbeat.exe devices
```

```
0: \Device\NPF_{644A5E7E-FFD4-4769-A2B0-5AB6601C98D2} (Intel(R)  
PRO/1000 MT Desktop Adapter) (fe80::9dfb:2aae:7112:2f1a  
192.168.1.103)
```

```
1: \Device\NPF_{D1F0606C-C08A-4F60-8DDE-4F840663AE98} (Intel(R)  
PRO/1000 MT Desktop Adapter) (fe80::b9b0:36e6:1851:225d  
10.0.4.15)
```

```
2: \Device\NPF_Loopback (Adapter for loopback traffic capture)  
(Not assigned ip address)
```

We're going to want to capture our network traffic on our connection to the internet (NAT). This is interface 1 – you can tell as it has the IP address of 10.0.4.15.

**Important note**

I have observed nuances with the lab environments and VirtualBox, where the captured device number can change. If you are not observing the data that you're expecting, re-run `packetbeat.exe` devices to ensure that the device number you've set in `packetbeat.yml` is still the interface with the IP address of 10.0.4.15.

With this administrative terminal, open `packetbeat.yml` and update the interface from 0 to 1:

```
# === Network device ===  
  
# Select the network interface to sniff the data. On Linux, you  
# can use the  
# "any" keyword to sniff on all connected interfaces.  
packetbeat.interfaces.device: 1
```

Continue to update this file with the same configuration settings used for Winlogbeat (the Kibana and Elasticsearch IP and credentials).

Let's test the configuration and output:

```
packetbeat.exe test config  
  
packetbeat.exe test output
```

Let's run the setup to load the index patterns and dashboards:

```
packetbeat.exe setup
```

Run Packetbeat as a binary and check on the **Discover** tab to see if you see the data (don't forget to switch to the `packetbeat-*` index pattern):

```
packetbeat.exe
```